

Movie Recommendation

August 3, 2024

1 BIG DATA CAPSTONE PROJECT : MILESTONE REPORT 4

2 PROFESSOR: MEYSAM EFFATI

3 TEAM: ARUN KUMAR SUBRAMANIAM, CARLOS MUNOZ EBRATT, KAMALDEEP KAUR, LUIS ALEJANDRO GUTIERREZ HAYEK, OSCAR FELIPE FERNANDEZ TOVAR

3.0.1 Importing the libraries

```
[1]: import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
```

3.0.2 Loading the dataset

```
[2]: movies_df = pd.read_csv('mflix.movies.csv')
```

3.0.3 Data Exploration

The `df.head()` function in pandas is used to display the first few rows of a DataFrame. By default, it shows the first 5 rows.

```
[3]: print(movies_df.head())
```

			plot	genres[0]	genres[1]	\
0	The cartoonist, Winsor McCay, brings the Dinos...		Animation		Short	
1	An immigrant leaves his sweetheart in Italy to...		Drama		NaN	
2	A rich young Easterner who has always wanted t...		Comedy		Western	
3	A penniless young man tries to save an heiress...		Comedy		Short	
4	A tipsy doctor encounters his patient sleepwal...		Comedy		Short	

	genres[2]	runtime	cast[0]	cast[1]	cast[2]	\
0	Comedy	12.0	Winsor McCay	George McManus	Roy L. McCardell	
1	NaN	78.0	George Beban	Clara Williams	J. Frank Burke	

2	Romance	72.0	Douglas Fairbanks	Eileen Percy	Calvert Carter
3	Action	22.0	Harold Lloyd	Mildred Davis	'Snub' Pollard
4	NaN	26.0	Harold Lloyd	Roy Brooks	Mildred Davis

	cast[3]	num_mflix_comments	...	awards.text	lastupdated	year	\
0	NaN	0	...	1 win.	03:15.3	1914	
1	Leo Willis	0	...	1 win.	07:43.2	1915	
2	Charles Stevens	0	...	1 win.	40:35.1	1917	
3	Peggy Cartwright	0	...	1 nomination.	16:14.2	1919	
4	Wallace Howe	1	...	1 nomination.	35:33.7	1920	

	imdb.rating	imdb.votes	imdb.id	countries[0]	type	tomatoes.viewer.rating	\
0	7.3	1837.0	4008	USA	movie	3.7	
1	6.4	175.0	5557	USA	movie	4.0	
2	6.9	388.0	8775	USA	movie	NaN	
3	7.0	639.0	10146	USA	movie	3.3	
4	7.0	646.0	11293	USA	movie	3.4	

	rated
0	NaN
1	PASSED
2	NaN
3	TV-G
4	PASSED

[5 rows x 31 columns]

The df.info() function in pandas provides a summary of a DataFrame df.

```
[4]: print(movies_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21349 entries, 0 to 21348
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   plot                                  20203 non-null  object
1   genres[0]                            21237 non-null  object
2   genres[1]                            15345 non-null  object
3   genres[2]                            8696 non-null   object
4   runtime                              20910 non-null  float64
5   cast[0]                              20987 non-null  object
6   cast[1]                              20686 non-null  object
7   cast[2]                              20484 non-null  object
8   cast[3]                              20309 non-null  object
9   num_mflix_comments                  21349 non-null  int64
10  poster                              18044 non-null  object
11  title                               21349 non-null  object
12  fullplot                            19852 non-null  object
```

```

13 languages[0]          21119 non-null object
14 released              20878 non-null object
15 directors[0]          21107 non-null object
16 directors[1]          1580 non-null object
17 writers[0]            20256 non-null object
18 writers[1]            13427 non-null object
19 awards.wins            21349 non-null int64
20 awards.nominations     21349 non-null int64
21 awards.text            21349 non-null object
22 lastupdated            21349 non-null object
23 year                   21349 non-null object
24 imdb.rating            21288 non-null float64
25 imdb.votes             21287 non-null float64
26 imdb.id                21349 non-null int64
27 countries[0]           21339 non-null object
28 type                   21349 non-null object
29 tomatoes.viewer.rating 18566 non-null float64
30 rated                  11455 non-null object
dtypes: float64(4), int64(4), object(23)
memory usage: 5.0+ MB
None

```

3.0.4 Data Handling

Getting the summary of missing values in each column

```

[5]: missing_values = movies_df.isnull().sum()
missing_values = missing_values[missing_values > 0]
missing_values

```

```

[5]: plot          1146
genres[0]          112
genres[1]          6004
genres[2]          12653
runtime            439
cast[0]            362
cast[1]            663
cast[2]            865
cast[3]            1040
poster            3305
fullplot           1497
languages[0]        230
released           471
directors[0]         242
directors[1]        19769
writers[0]           1093
writers[1]           7922
imdb.rating          61

```

```
imdb.votes          62
countries[0]        10
tomatoes.viewer.rating 2783
rated              9894
dtype: int64
```

Dropping columns with excessive missing values

```
[6]: columns_to_drop = ['directors[1]', 'writers[1]', 'genres[2]']
     movies_df.drop(columns=columns_to_drop, inplace=True)
```

Merging genres and Cast into a single column

```
[7]: movies_df['Genre'] = movies_df[['genres[0]', 'genres[1]']].apply(lambda x: ', '.join(x.dropna().astype(str)), axis=1)
```

```
[8]: movies_df['Cast'] = movies_df[['cast[0]', 'cast[1]', 'cast[2]', 'cast[3]']].apply(lambda x: ', '.join(x.dropna().astype(str)), axis=1)
```

Dropping the columns that were merged

```
[9]: movies_df.drop(columns=['genres[0]', 'genres[1]', 'cast[0]', 'cast[1]', 'cast[2]', 'cast[3]'], inplace=True)
```

Removing rows with null values in specific columns because we can't fill the missing values with mode or media, since these are too unique.

```
[10]: columns_to_check = ['poster', 'fullplot', 'languages[0]', 'released', 'directors[0]', 'writers[0]', 'countries[0]', 'poster', 'awards.text', 'lastupdated', 'imdb.votes', 'type', 'rated']
     movies_df.dropna(subset=columns_to_check, inplace=True)
```

```
[11]: numerical_columns = ['runtime', 'imdb.rating', 'imdb.votes', 'tomatoes.viewer.rating']
     for column in numerical_columns:
         movies_df[column].fillna(movies_df[column].median(), inplace=True)
```

Filling missing values in 'rated' with a placeholder

```
[12]: movies_df['rated'].fillna('Not Rated', inplace=True)
```

Checking for missing values

```
[13]: missing_values = movies_df.isnull().sum()
     missing_values = missing_values[missing_values > 0]
     missing_values
```

```
[13]: Series([], dtype: int64)
```

Renaming the columns with understandable names.

```
[14]: movies_df.rename(columns={
    'plot': 'Plot',
    'runtime': 'Runtime',
    'num_mflix_comments': 'Comments',
    'title': 'Title',
    'fullplot': 'FullPlot',
    'languages[0]': 'Language',
    'released': 'Released',
    'directors[0]': 'Director',
    'writers[0]': 'Writer',
    'awards.wins': 'Awards',
    'awards.nominations': 'Nominations',
    'year': 'Year',
    'imdb.rating': 'IMDB Rating',
    'countries[0]': 'Countries',
    'tomatoes.viewer.rating': 'Tomatoes Rating'
}, inplace=True)
```

Now checking the dataframe information

```
[15]: # Display the updated dataset information and first few rows
updated_info = movies_df.info()
updated_head = movies_df.head()
updated_info, updated_head
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 10701 entries, 3 to 21317
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Plot                  10701 non-null  object
1   Runtime               10701 non-null  float64
2   Comments              10701 non-null  int64
3   poster                10701 non-null  object
4   Title                 10701 non-null  object
5   FullPlot              10701 non-null  object
6   Language              10701 non-null  object
7   Released              10701 non-null  object
8   Director              10701 non-null  object
9   Writer                10701 non-null  object
10  Awards                10701 non-null  int64
11  Nominations           10701 non-null  int64
12  awards.text           10701 non-null  object
13  lastupdated           10701 non-null  object
14  Year                  10701 non-null  object
15  IMDB Rating           10701 non-null  float64
16  imdb.votes            10701 non-null  float64
17  imdb.id               10701 non-null  int64
```

```

18 Countries          10701 non-null object
19 type                10701 non-null object
20 Tomatoes Rating    10701 non-null float64
21 rated              10701 non-null object
22 Genre              10701 non-null object
23 Cast               10701 non-null object
dtypes: float64(4), int64(4), object(16)
memory usage: 2.0+ MB

```

[15]: (None,

```

Plot Runtime Comments \
3 A penniless young man tries to save an heiress... 22.0 0
4 A tipsy doctor encounters his patient sleepwal... 26.0 1
5 A young man, unaccustomed to children, must ac... 35.0 0
9 Millie Stope lives with her grandfather on a r... 88.0 1
11 Mrs Erlynne, the mother of Lady Windermere - h... 120.0 1

```

```

poster Title \
3 https://m.media-amazon.com/images/M/MV5BNzE10W... From Hand to Mouth
4 https://m.media-amazon.com/images/M/MV5BODliMj... High and Dizzy
5 https://m.media-amazon.com/images/M/MV5BYjgzYz... Now or Never
9 https://m.media-amazon.com/images/M/MV5BMjA40T... Wild Oranges
11 https://m.media-amazon.com/images/M/MV5BZDUyYz... Lady Windermere's Fan

```

```

FullPlot Language \
3 As a penniless man worries about how he will m... English
4 After a long wait, a young doctor finally has ... English
5 Mary is looking after a young child whose pare... English
9 Millie Stope lives with her grandfather on a r... English
11 Mrs Erlynne, the mother of Lady Windermere - h... English

```

```

Released Director Writer \
3 1919-12-28T00:00:00.000Z Alfred J. Goulding H.M. Walker (titles)
4 1920-07-11T00:00:00.000Z Hal Roach Frank Terry (story)
5 1921-03-27T00:00:00.000Z Fred C. Newmeyer H.M. Walker (titles)
9 1924-01-20T00:00:00.000Z King Vidor Joseph Hergesheimer (by)
11 1925-12-26T00:00:00.000Z Ernst Lubitsch Oscar Wilde (by)

```

```

... Year IMDB Rating imdb.votes imdb.id Countries type \
3 ... 1919 7.0 639.0 10146 USA movie
4 ... 1920 7.0 646.0 11293 USA movie
5 ... 1921 6.8 489.0 12512 USA movie
9 ... 1924 7.1 327.0 15498 USA movie
11 ... 1925 7.6 630.0 16004 USA movie

```

```

Tomatoes Rating rated Genre \
3 3.3 TV-G Comedy, Short

```

4	3.4	PASSED	Comedy, Short
5	3.8	PASSED	Comedy, Short
9	4.2	PASSED	Drama, Romance
11	3.7	APPROVED	Comedy

Cast

3	Harold Lloyd, Mildred Davis, 'Snub' Pollard, P...
4	Harold Lloyd, Roy Brooks, Mildred Davis, Walla...
5	Harold Lloyd, Mildred Davis, Anna Mae Bilson
9	Frank Mayo, Virginia Valli, Ford Sterling, Nig...
11	Ronald Colman, May McAvoy, Bert Lytell, Irene ...

[5 rows x 24 columns])

3.0.5 Model 1

Vectorize the Genre column using TF-IDF

```
[16]: tfidf_vectorizer = TfidfVectorizer()
      tfidf_matrix = tfidf_vectorizer.fit_transform(movies_df['Genre'])
```

Calculate the cosine similarity between the genre vectors

```
[17]: cosine_sim = cosine_similarity(tfidf_matrix, tfidf_matrix)
```

Function to get movie recommendations based on genre similarity

```
[18]: def get_recommendations(title, cosine_sim=cosine_sim):
      # Get the index of the movie that matches the title
      idx = movies_df[movies_df['Title'].str.lower() == title.lower()].index[0]

      # Get the pairwise similarity scores of all movies with that movie
      sim_scores = list(enumerate(cosine_sim[idx]))

      # Sort the movies based on the similarity scores
      sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

      # Get the indices of the 10 most similar movies
      sim_scores = sim_scores[1:11] # Exclude the first movie (itself)

      # Get the movie indices
      movie_indices = [i[0] for i in sim_scores]

      # Return the top 10 most similar movies
      return movies_df['Title'].iloc[movie_indices]
```

Example usage: Get recommendations for a specific movie

```
[19]: recommendations = get_recommendations("A Walk in the Sun")
recommendations
```

```
[19]: 95          Eskimo
113    The Emperor Jones
117          Comradeship
147    Dante's Inferno
155    Come and Get It
182    The Informer
269    Of Mice and Men
280    The Green Pastures
302    The Citadel
311          Tevya
Name: Title, dtype: object
```

3.0.6 Model 2

Function to get movie recommendations based on genre and released year as input

Combine Genre and Released Year into a new feature

```
[20]: movies_df['Genre_Released'] = movies_df['Genre'] + ' ' + movies_df['Released'].
      ↪astype(str)
```

Vectorize the combined Genre_Released column using TF-IDF

```
[21]: tfidf_vectorizer_combined = TfidfVectorizer(max_features=500)
tfidf_matrix_combined = tfidf_vectorizer_combined.
      ↪fit_transform(movies_df['Genre_Released'])
```

```
[22]: def recommend_movies_by_genre_year(genre, year,
      ↪tfidf_matrix=tfidf_matrix_combined, n_recommendations=10):
    # Combine genre and year into a single string
    input_combined = genre + ' ' + str(year)

    # Transform the input into a TF-IDF vector
    input_vector = tfidf_vectorizer_combined.transform([input_combined])

    # Calculate the cosine similarity of the input with all movies
    cosine_similarities = cosine_similarity(input_vector,
      ↪tfidf_matrix_combined).flatten()

    # Get the indices of the most similar movies
    sim_indices = cosine_similarities.argsort()[-(n_recommendations + 1):][::
      ↪-1][1:]

    # Return the top n most similar movies
    return movies_df['Title'].iloc[sim_indices]
```


Example usage: Get recommendations for a specific genre and release year

```
[23]: recommended_movies = recommend_movies_by_genre_year("Action", 2012)
recommended_movies
```

```
[23]: 15755          Bellflower
15815    The Man with the Iron Fists
17779          The Day
17210          My Way
17745          Here Comes the Boom
17362          Here Comes the Boom
16438          Men in Black 3
17636          Lockout
17179          Lockout
18479          Tai Chi Zero
Name: Title, dtype: object
```

3.0.7 Movie Recommendation System: Model Report

Introduction This project aims to build a movie recommendation system that provides movie suggestions based on the genre and release year input by a user. To achieve this, we have utilized a content-based filtering approach using Term Frequency-Inverse Document Frequency (TF-IDF) vectorization and cosine similarity.

Models and Techniques Used

1. TF-IDF Vectorization:

- **Purpose:** To convert textual data (genres and release years) into numerical vectors that can be used for similarity calculations.
- **Why TF-IDF:** TF-IDF is a statistical measure used to evaluate the importance of a word in a document relative to a collection of documents (corpus). It helps in transforming the textual data into a matrix of features, capturing the relevance of terms in each movie description.
- **Implementation:** The `TfidfVectorizer` from the `sklearn.feature_extraction.text` module was used with a maximum feature limit to manage memory usage.

2. Cosine Similarity:

- **Purpose:** To calculate the similarity between the input movie (based on genre and release year) and all other movies in the dataset.
- **Why Cosine Similarity:** Cosine similarity measures the cosine of the angle between two vectors, providing a measure of how similar the two vectors are irrespective of their magnitude. It is particularly useful for text data, as it helps to find the direction of the vectors (representing documents) rather than their lengths.
- **Implementation:** The `cosine_similarity` function from the `sklearn.metrics.pairwise` module was used to compute the similarity between the TF-IDF vectors.

Steps in the Model

1. Data Preparation:

- Combined the genre and release year of each movie into a single feature (**Genre_Released**) to provide a comprehensive representation of both attributes.
- Example: “Drama 2000” combines the genre “Drama” with the release year “2000”.

2. Vectorization:

- Applied TF-IDF vectorization to the **Genre_Released** feature to convert the text data into a numerical matrix (**tfidf_matrix_combined**).

3. Similarity Calculation:

- For a given input (genre and release year), transformed the input into a TF-IDF vector.
- Calculated the cosine similarity between the input vector and all other movie vectors in the dataset.

4. Recommendation Generation:

- Sorted the movies based on similarity scores and selected the top 10 most similar movies
= `recommend_movies_by_genre_year(“Drama”, 2000)` ““

Conclusion The content-based filtering approach using TF-IDF vectorization and cosine similarity effectively provides movie recommendations based on the genre and release year. This method leverages textual features to understand the context and similarity between movies, making it a robust solution for personalized movie recommendations.